

CIS 5800

# Machine Perception

Instructor: Lingjie Liu

Lec 13: March 5, 2026

# Administrivia

- Next week is Spring Break
- Deadlines of HW 2 and HW3 are March 17 and March 18 respectively
- Midterm Exam (March 26)
  - The format is slightly different from that of the review questions, but the coverage of the key knowledge points will be the same.
  - 5 long questions (50pts in total = 10pts x 5)
  - Knowledge Points: Everything covered up to the class of March 17.
- Combined slides for your mid-term exam preparation will be uploaded to Canvas tomorrow.
- Review for Knowledge Points: Next Class (March 24).
- We will have additional OHs in the week after Spring Break. TAs will announce the specific times asap.

## 2-View SfM: How we recovered $R, T, \{\lambda_i, \mu_i\}_{i=1}^n$


- Invent an object  $E = \hat{T}R$
- Express epipolar constraint for every point correspondence  $q_i^T E p_i = 0$
- Solve equations to get initial estimate  $E_0$
- $SVD(E_0) = U\Sigma V^T$ , Then set  $E = U\text{diag}(1, 1, 0)V^T$
- Then, long proof to show you can set  $\pm$ (last left singular vector of  $E$ ) to be  $T$ , and there are two possible rotation matrices you can pair each translation vector with.
  - Solve depth by solving  $\mu_i R p_i + T = \lambda_i q_i$  for each of these 4 solutions.
  - Settle on the solution that has all (or most) depths positive.  
(Sidenote/Q: Are  $\lambda_i$  and  $\mu_i$  really “depths?” )

**What is the meaning of these ambiguities in the solution to SfM?**

# Types of Ambiguity in 2-View SfM Solutions


**Mirror ambiguity:** If  $T$  is a solution, then  $-T$  is a solution, too. There is no way to disambiguate from the epipolar constraint:  $q^T(-T \times Rp) = 0$ .

Just extending scale ambiguity to negative scales. Some people don't even count this as a "true" ambiguity of SfM.




**Twisted pair ambiguity:** If  $R$  is a solution, then also  $R_{T,\pi}R$  is a solution. The first camera is "twisted" around the baseline 180 degrees.


# Mirror ambiguity




(a)




(b)




(c)



(d)



Twisted pair ambiguity



# What *is* “solving” when solutions have “ambiguity”?

- We said we could solve for camera pose given 3 2D->3D correspondences (P3P)
  - But there was ambiguity: 4 solutions!

- We said

- But t  
into  
of E)

## P3P Step 1: The algebraic drudgery

Reduces to two quadratic equations in  $u$  and  $v$ .

$$d_{13}^2(u^2 + v^2 - 2uv \cos \delta_{23}) = d_{23}^2(1 + v^2 - 2v \cos \delta_{13}) \quad (1)$$

$$d_{12}^2(1 + v^2 - 2v \cos \delta_{13}) = d_{13}^2(u^2 + 1 - 2u \cos \delta_{12}) \quad (2)$$

- So what

- We i

- In ot  
spac

- We i  
pin c

- This  
setu  
and  
prok

- Solve Eqn (1) for  $u^2$  in terms of  $u, v, v^2$  (and constants).
- Plug this solution into Eqn (2), so that it has no  $u^2$  term. Solve for  $u$  in terms of  $v, v^2$ , and constants.
- Plug this solution for  $u$  back into Eqn (1), so that it has no more  $u$  or  $u^2$ . Instead, it is a 4<sup>th</sup> degree equation in  $v$ . *Get the 4 real solutions analytically.*
- Then plug back into the solution found in step b) above, to get  $u$ .
- Then get  $d_1$  from the quadratic equations on the last page.
- Then plug back into  $d_2 = ud_1$  and  $d_3 = vd_1$  from the last page to get  $d_2, d_3$ .

ch E  
ty

) to

em  
ms,  
xed






Recap: What is the relationship between two views of the same plane? (facade)



Answer: As we saw early in the course, homographies!



# Recap: Two Calibrated Views of the Same 3D Scene



$$R(\lambda p) + T = \mu q$$

Given 2D correspondences  $(p, q)$

Find motion  $R, T$  and depths  $\lambda, \mu$ .


# Recap: In 2-view SfM, what if there is no translation?

- Epipolar constraint becomes  $0=0$ ! So, cannot do all the SfM stuff above!
- Go back to  $\lambda q = \mu Rp + T(= 0) = \mu Rp$ 
  - In other words:  $q \sim Rp$
  - Going back to pixel coordinates from calibrated:
    - $K^{-1}q_{px} \sim RK^{-1}p_{px}$
    - $q_{px} \sim KRK^{-1}p_{px}$
- $KRK^{-1}$  is invertible (determinant  $\neq 0$ )
- So this is a homography  $H = KRK^{-1}$ !

**Two views from the same camera center and different camera orientations are related by a homography even if the world is not planar!**

# No-Translation Image = Image of a Plane!

**Two views from the same camera center and different camera orientations are related by a homography even if the world is not planar!**



The image formed of the world on plane 2, may also be thought of as:  
the image of image plane 1. i.e. image of a plane i.e. homography!

# Recap: Checking for no translation during SfM

- If you set up the  $n$  point correspondences as  $A_{2n \times 9} \mathbf{h}_{9 \times 1} = 0$  (as we have done before when solving for homography)
- Then, if  $\text{rank}(A)$  is (approximately) 8, then you can (approximately) compute  $H$ . This means that, actually, there was no (or insignificant) translation, so you can't do SfM!

Q: Is this the only case when  $H$  is computable?

A: No, also when imaging a plane, of course.

This check is a common component in SfM systems.

Note: While you can't do *SfM* with no translation of the camera, you can actually do other cool things, like building 360° image “panoramas”!



Coming up next!





# Image Stitching / Mosaicing From Rotated Views

# 2 Views From The Same Camera Center



- Camera field of view (FOV) is finite, so view 1 can only fill in pixels into a small region of image plane 1
- View 2 could see things outside the FOV of view 1. Could we add those pixels into view 1 to extend the image?

**Yes. Project both images onto the same image plane e.g. image 1!**


Your human vision system does something like this when your eyeball moves.

# Image Stitching Results Example



# What is the transformation between these views?

For every pixel in image 2, we need to find the right pixel location to place it inside image 1.



**3x3 homography**

$$\begin{bmatrix} x_1 \\ y_1 \\ 1 \end{bmatrix} \sim \mathbf{K}_1 \mathbf{R}_1 \mathbf{R}_2^T \mathbf{K}_2^{-1} \begin{bmatrix} x_2 \\ y_2 \\ 1 \end{bmatrix}$$

image coords (in image 1)      image coords (in image 2)

3D ray direction (in cam 2 coordinates)

3D ray direction in world coordinates

3D ray direction in camera 1 coordinates

Project through camera 1!

Homography even though not restricting to imaging a plane!



# In practice, the homography is *estimated*

- We now know that the two images are related by a homography.
- In practice, we don't know  $R$  or even necessarily  $K_1$  or  $K_2$
- But fortunately, we know how to solve for homography  $H$  given point correspondences.
  - So we need the two images to overlap so that you can find correspondences and estimate the homography.

# Creating a panorama from >2 images

- Basic Procedure
  - Take a sequence of images from the same position
    - Rotate the camera about its optical center
  - Compute transformation between second image and first
  - Transform the second image to overlap with the first
  - Blend the two together to create a mosaic
  - If there are more images, repeat


# Projecting images onto a common plane





# Image alignment

Often cropped like this to produce an image without empty regions






# Can we use homography to create a 360° panorama?



Microsoft Lobby: <http://www.acm.org/pubs/citations/proceedings/graph/258734/p251-szeliski>

# Can we use homography to create a 360° panorama?




Projecting backwards-facing views onto the same image plane?  
Breaks down.


# Panoramas

**Instead of projecting onto a plane, you project onto a sphere!**

i.e. for every 3D point, draw the line connecting it to the camera center, and find its intersection with a projection sphere --- this is the point's image!



# Spherical projection (overview)




- Map 3D point  $(X, Y, Z)$  onto sphere

$$(\hat{x}, \hat{y}, \hat{z}) = \frac{1}{\sqrt{X^2 + Y^2 + Z^2}}(X, Y, Z)$$

- Convert to spherical coordinates  
 $(\sin\theta\cos\phi, \sin\phi, \cos\theta\cos\phi) = (\hat{x}, \hat{y}, \hat{z})$
- Convert to spherical image coordinates

$$(\tilde{x}, \tilde{y}) = (s\theta, s\phi) + (\tilde{x}_c, \tilde{y}_c)$$

- $s$  defines size of the final image



# Spherical Panorama Example



Note the  
distortions

Microsoft Lobby: <http://www.acm.org/pubs/citations/proceedings/graph/258734/p251-szeliski>





# Hough and RANSAC



Simple subroutines for feature identification, grouping, and correspondence filtering



*And now for something completely different*

# Correspondences are not clean!


Or we might have to group them !  
(two planes=>two homographies)



# A Case Study: Fitting A Line To Data

Given data  $(x_i, y_i)$  belonging to a line


$$x \cos \theta + y \sin \theta = d$$



# Finding the line by optimization

To solve for the line, we could minimize the mean squared error:

$$\operatorname{argmin}_{\theta \in [0, 2\pi), d \geq 0} \sum_{i=1}^N (x_i \cos \theta + y_i \sin \theta - d)^2$$



This is solvable. Exactly how is not particularly important. (next 2 slides)

# Solution sketch (not for testing)

$$\operatorname{argmin}_{\theta \in [0, 2\pi), d \geq 0} \sum_{i=1}^N (x_i \cos \theta + y_i \sin \theta - d)^2$$

$$f(\theta, d) = \sum_{i=1}^N (x_i \cos \theta + y_i \sin \theta - d)^2$$

$$\frac{\partial f}{\partial d} = \sum_{i=1}^N (-2)(x_i \cos \theta + y_i \sin \theta - d) = 0$$

$$d = \frac{1}{N} \sum_i (x_i \cos \theta + y_i \sin \theta)$$

# Solution sketch (not for testing)

We can eliminate  $d$  and obtain

$$\arg \min_{\theta \in [0, 2\pi)} \sum_{i=1}^N ((x_i - \bar{x}) \cos \theta + (y_i - \bar{y}) \sin \theta)^2$$

which is equivalent to the minimization of the Rayleigh quotient

$$\arg \min_{\theta \in [0, 2\pi)} \begin{pmatrix} \cos \theta & \sin \theta \end{pmatrix} \underbrace{\sum_{i=1}^N \begin{pmatrix} (x_i - \bar{x})^2 & (x_i - \bar{x})(y_i - \bar{y}) \\ (x_i - \bar{x})(y_i - \bar{y}) & (y_i - \bar{y})^2 \end{pmatrix}}_C \begin{pmatrix} \cos \theta \\ \sin \theta \end{pmatrix}$$

Solution is the eigenvector to the minimal eigenvalue of  $C$ .

This corresponds to maximizing:


$$\sum_i e^{-\frac{1}{2\sigma_i^2}(\cos \theta x_i + \sin \theta y_i - d)^2}$$

with  $\sigma_i = 1$  being the same for all points.

$P(\theta, d|x_i, y_i)$ , kind of like a “vote” by  $(x_i, y_i)$  for  $(\theta, d)$

If you have machine learning or statistics background, what is this type of line fitting called?

Maximum likelihood!



Key idea: find the parameters that the most data points “vote” for.



# Line Parameter Space

For a line in the form:

$$x \cos \theta + y \sin \theta = d$$

The parameter space of the line is in  $(\theta, d)$ . Each point in this parameter space corresponds to a line in 2D.


For a point  $(x_i, y_i)$  in x-y space, what does it correspond to in  $(\theta, d)$ ?

$$d = x_i \cos \theta + y_i \sin \theta$$

Line through two points  $(x_1, y_1)$  and  $(x_2, y_2)$  can be found as the intersection  $(\theta, d)$  of the two curves:

$$d = \cos \theta x_1 + \sin \theta y_1$$

$$d = \cos \theta x_2 + \sin \theta y_2$$



# Hough Voting Procedure

For every possible line  $(d, \theta)$  , count the number of points that “support” it.


$$\sum_{\text{points } x,y} 1(x \cos \theta - y \sin \theta = d)$$

(can sometimes relax the equality sign to approximate equality, or makes vote “soft” etc.)

In the end, the  $(d, \theta)$  with the maximum “votes” wins!

(in case of multiple lines, not just one winner, but all local maxima above some threshold win)

# Vote Counting to Find Lines!



# Algorithm v1, and a practical issue

Seems impractical

- For every possible line  $(d, \theta)$ , count the number of points that “support” it.

$$\sum_{\text{points } x,y} 1(|x \cos \theta + y \sin \theta - d| < \delta_{\text{threshold}})$$

- In the end, local maxima with many votes are declared lines.

# Solution: discretizing the parameter space.

- For each of some **enumerated discrete parameter combinations**  $\{(d_j, \theta_j)\}_{j=1}^J$ , count supporting points among the data.

$$V[d_j, \theta_j] = \sum_{\text{points } x,y} 1(|x \cos \theta_j + y \sin \theta_j - d_j| < \delta_{\text{threshold}})$$

- For example,  $\Theta = \{0^\circ, 15^\circ, \dots, 180^\circ\} \times D = \{0, 1, 2, 3, 4, 5\} \Rightarrow 13 \times 6 = 78$  parameter combinations
- In the end, **local maxima of  $V[d, \theta]$**  with many votes are declared lines.



## Notes:

Needs careful choices about how to discretize parameters into bins.

Each bin should correspond to an error you are okay with making, so can't be too coarse.

But too fine-grained => computationally intensive!

# Line Fitting Demo With $>1$ Lines




# Hough Circle Detection (3 parameters)

If we fix the  $(x, y)$  in the circle equation

$$(x - a)^2 + (y - b)^2 = r$$

we obtain the equation of a cone in  $(a, b, r)$  Hough space. Observe that we used  $r$  and not  $r^2$ .






# Hough Ellipse Detection (5 parameters)

Let us look at the ellipse equation

$$(x - c_x \quad y - c_y) \begin{pmatrix} \cos \theta & \sin \theta \\ -\sin \theta & \cos \theta \end{pmatrix} \begin{pmatrix} \frac{1}{\alpha^2} & 0 \\ 0 & \frac{1}{\beta^2} \end{pmatrix} \begin{pmatrix} \cos \theta & -\sin \theta \\ \sin \theta & \cos \theta \end{pmatrix} \begin{pmatrix} x - c_x \\ y - c_y \end{pmatrix} = 1$$



Hough transform needs the computation of  $V(c_x, c_y, \alpha, \beta, \theta)$  and search for maxima in 5-dimensional space. Hough becomes intractable beyond lines and circles.

# Application: Automatically Grouping Vanishing Points

(Roughly, each pair of detected line segments votes for their point of intersection. Larger weights for longer segments, and non-collinear segments)



# Disadvantages of Hough Transforms

Hough transforms are great and simple-to-implement approach for grouping  $k$  samples into  $n$  “groups” (e.g. lines, circles, VPs etc.) + outliers that don’t belong to any group.

- Note: Needs a bounded parameter space to allow a finite discrete set of hypotheses, and careful discretization.
- **Main drawback:** Poor scaling. Becomes intractable when # parameters(unknowns) for the target groupings is more than 3 or at most 4 in practice.
  - So, e.g. to fit homographies (8 parameters), completely intractable.

Alternative solution that scales much better: RANSAC!

